

K nearest neighbor

Fatemeh Mansoori

This slides are created based on the slides of the Sebastian Raschka for the introduction to machine learning course



Topics

- **Intro to nearest neighbor models**
- Nearest neighbor decision boundary
- K-nearest neighbors
- Improving k-nearest neighbors: modifications and hyperparameters
- K-nearest neighbors in Python

1-Nearest Neighbor

1-Nearest Neighbor

- Task: predict the target / label of a new data point



- How? Look at most "similar" data point in training set

1-Nearest Neighbor Training Step

$$\langle \mathbf{x}^{[i]}, y^{[i]} \rangle \in \mathcal{D} \quad (|\mathcal{D}| = n)$$

How do we "train" the 1-NN model?

To train the 1-NN model, we simply "remember" the training dataset

1-Nearest Neighbor Prediction Step

closest_point := None

closest_distance := ∞

- for $i = 1, \dots, n$:
 - current_distance := $d(\mathbf{x}^{[i]}, \mathbf{x}^{[q]})$
 - if current_distance < closest_distance:
 - closest_distance := current_distance
 - closest_point := $\mathbf{x}^{[i]}$
- return $f(\text{closest_point})$

query point

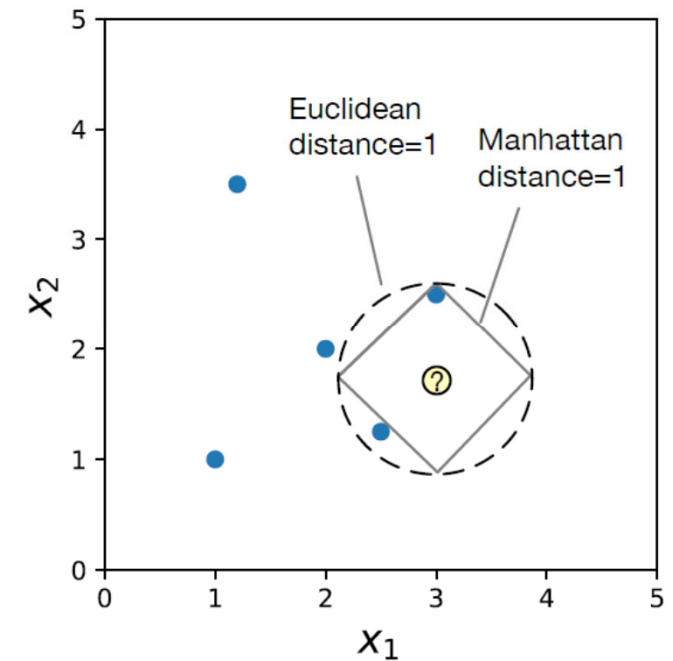


Which Point is Closest to ? ?

- Depends on the Distance Measure!
- Commonly used: Euclidean Distance

$$d(\mathbf{x}^{[a]}, \mathbf{x}^{[b]}) = \sqrt{\sum_{j=1}^m \left(x_j^{[a]} - x_j^{[b]} \right)^2}$$

- Other metrics : Manhattan distance
-



Some Common Continuous Distance Measures

- Euclidean

- Manhattan

- Minkowski: $d(\mathbf{x}^{[a]}, \mathbf{x}^{[b]}) = \left[\sum_{j=1}^m \left(\left| x_j^{[a]} - x_j^{[b]} \right| \right)^p \right]^{\frac{1}{p}}$

- Cosine similarity

Some Discrete Distance Measures



Hamming distance: $d(\mathbf{x}^{[a]}, \mathbf{x}^{[b]}) = \sum_{j=1}^m \left| x_j^{[a]} - x_j^{[b]} \right|$ where $x_j \in \{0,1\}$

Jaccard/Tanimoto similarity:

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$$


Dice: $D(A, B) = \frac{2|A \cap B|}{|A| + |B|}$

Feature Scaling





Nearest Neighbor Decision Boundary

Decision Boundary Between (a) and (b)




Decision Boundary of 1-NN

Using Delaunay triangulation



K nearest neighbour

k-Nearest Neighbors



A

y: ● ● ● ● ■ ■ ■ ■ ■ ■

Majority vote: ■

Plurality Vote: ■

B

y: ● ● ● ■ ■ ■ ◆ ◆ ◆ ◆

Majority vote: None

Plurality Vote: ◆

kNN for Classification

$$\mathcal{D}_k = \{ \langle \mathbf{x}^{[1]}, f(\mathbf{x}^{[1]}) \rangle, \dots, \langle \mathbf{x}^{[k]}, f(\mathbf{x}^{[k]}) \rangle \} \quad \mathcal{D}_k \subseteq \mathcal{D}$$

$$h(\mathbf{x}^{[q]}) = \underset{y \in \{1, \dots, t\}}{\operatorname{arg\,max}} \sum_{i=1}^k \delta(y, f(\mathbf{x}^{[i]}))$$
$$\delta(a, b) = \begin{cases} 1, & \text{if } a = b, \\ 0, & \text{if } a \neq b. \end{cases}$$

$$h(\mathbf{x}^{[t]}) = \mathbf{mode}(\{f(\mathbf{x}^{[1]}), \dots, f(\mathbf{x}^{[k]})\})$$

*k*NN for Regression

$$\mathcal{D}_k = \{ \langle \mathbf{x}^{[1]}, f(\mathbf{x}^{[1]}) \rangle, \dots, \langle \mathbf{x}^{[k]}, f(\mathbf{x}^{[k]}) \rangle \} \quad \mathcal{D}_k \subseteq \mathcal{D}$$

$$h(\mathbf{x}^{[t]}) = \frac{1}{k} \sum_{i=1}^k f(\mathbf{x}^{[i]})$$

Distance-weighted k NN

$$h(\mathbf{x}^{[t]}) = \arg \max_{j \in \{1, \dots, p\}} \sum_{i=1}^k w^{[i]} \delta(j, f(\mathbf{x}^{[i]}))$$

$$w^{[i]} = \frac{1}{d(\mathbf{x}^{[i]}, \mathbf{x}^{[t]})^2}$$

Nearest Neighbor Search


$\mathcal{D}_k := \{\}$

while $|\mathcal{D}_k| < k$:

- `closest_distance := ∞`
- for $i = 1, \dots, n$, $\forall i \notin \mathcal{D}_k$:
 - `current_distance := $d(\mathbf{x}^{[i]}, \mathbf{x}^{[q]})$`
 - if `current_distance < closest_distance`:
 - * `closest_distance := current_distance`
 - * `closest_point := $\mathbf{x}^{[i]}$`
- add `closest_point` to \mathcal{D}_k

Improving Computational Performance

- Pruning




Hyperparameter


Value of k

$$k \in \{1, 3, 7\}$$


$k = _$



$k = _$



$k = _$



How to choose the value of K

- **Cross-Validation**

- try different values of K and evaluate the performance of the model using metrics like accuracy, precision, recall, or F1 score

- **Odd vs. Even K**

- It is generally recommended to use an odd value for K to avoid ties in voting

- **Rule of Thumb**

- A common rule of thumb is to set K to the square root of the total number of samples in your dataset. This is a good starting point but may not always be the optimal choice.

How to choose the value of K

- **Domain Knowledge**

- Depending on the characteristics of your dataset and problem domain, you may have insights that can guide you in choosing an appropriate value of K . For example, if you know that the decision boundaries are complex, you may want to choose a smaller value of K .

```
[ ] from sklearn.datasets import load_iris
    from sklearn.model_selection import train_test_split
    from sklearn.neighbors import KNeighborsClassifier
```

```
[ ] iris = load_iris()
    X, y = iris.data[:, 2:], iris.target
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                         shuffle=True)
```

```
[ ] knn_model = KNeighborsClassifier(n_neighbors=3)
    knn_model.fit(X_train, y_train)

    y_pred = knn_model.predict(X_test)
```

```
[ ] num_correct_predictions = (y_pred == y_test).sum()
    accuracy = (num_correct_predictions / y_test.shape[0]) * 100

    # print('Test set accuracy: %.2f%%' % accuracy)

    print(f'Test set accuracy: {accuracy:.2f}%')
```

Test set accuracy: 95.56%

```

mean_scores = []
for k in range (1,11) :
    knn_model = KNeighborsClassifier(n_neighbors = k)
    knn_model.fit(X_train, y_train)

    scores = cross_val_score(knn_model, X_train, y_train, cv = 5)
    mean_scores.append(scores.mean())

ck = np.argmax(mean_scores)
print(ck)
knn_model = KNeighborsClassifier(n_neighbors = ck+1)
knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_test)

num_correct_predictions = (y_pred == y_test).sum()
accuracy = (num_correct_predictions / y_test.shape[0]) * 100

# print('Test set accuracy: %.2f%%' % accuracy)

print(f'Test set accuracy: {accuracy:.2f}%')

```

4

Test set accuracy: 97.78%