


Linear Regression

What is linear ?

- Remember this:
 - $Y=mX+B$

A slope of 2 means that every 1-unit change in X yields a 2-unit change in Y.




Linear regression example

- Housing prices prediction

$$h_{\theta}(x) = \theta_0 + \theta_1 \times \text{frontage} + \theta_2 \times \text{depth}$$

- Estimating body weight by person height
- Estimate the happiness score of a person using its income



Formulation

- We generally formulate the linear regression model in matrix form:
 - $Y = Xw + \varepsilon$
- the target value y_i can be evaluated by
 - $y_i = \theta_0 + \theta_1 x_{i1} + \dots + \theta_n x_{in} + \varepsilon_i$
 - Y represents a vector of length n containing the observed values
 $Y = (y_1, \dots, y_m)^T$
 - ε is a vector for errors $\varepsilon = (\varepsilon_1, \dots, \varepsilon_m)^T$
 - X is a matrix of the features in which the column of ones incorporate the intercept


$$X = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1n} \\ 1 & x_{21} & x_{22} & \dots & x_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m1} & x_{m2} & \dots & x_{mn} \end{pmatrix}$$

- Hypothesis:

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d = \sum_{j=0}^d \theta_j x_j$$

Assume $x_0 = 1$

- Fit model by minimizing sum of squared errors



least squares (LSQ)



The fitted line is used as a predictor

Least square linear regression

- Cost Function

$$J(\boldsymbol{\theta}) = \frac{1}{2n} \sum_{i=1}^n \left(h_{\boldsymbol{\theta}} \left(\mathbf{x}^{(i)} \right) - y^{(i)} \right)^2$$

- Fit by solving $\min_{\boldsymbol{\theta}} J(\boldsymbol{\theta})$




Optimization Problem

One suitable estimator of β should be the one minimizing the sum of the squared errors

$$\|\epsilon\|_2^2 = \sum_{i=1}^m \epsilon_i^2 = \epsilon^T \epsilon.$$

$$\begin{aligned} \bullet \sum_{i=1}^m \epsilon_i^2 &= \epsilon^T \epsilon = (Y - X\theta)^T (Y - X\theta) \\ &= Y^T Y - 2\theta^T X^T Y + \theta^T X^T X \theta \end{aligned}$$



Vector derivative

$f(\mathbf{x})$	\rightarrow	$\frac{df}{d\mathbf{x}}$
$\mathbf{x}^T \mathbf{B}$	\rightarrow	\mathbf{B}
$\mathbf{x}^T \mathbf{b}$	\rightarrow	\mathbf{b}
$\mathbf{x}^T \mathbf{x}$	\rightarrow	$2\mathbf{x}$
$\mathbf{B} \mathbf{x}$	\rightarrow	\mathbf{B}^T

- Differentiating this term and setting it to zero, we find that the estimate for θ , which minimizes the squared error, satisfies the equation :

$$X^T X \theta = X^T Y$$


- Provided $X^T X$ is invertible :

$$\theta = (X^T X)^{-1} X^T Y$$

Geometric Approach


Another way to looking at this problem is to say we want a solution that lies in the space spanned by X become as close as possible to Y .

In this way, the systematic component $X\theta$ is projection of Y onto space spanned by X and residuals are $Y - X\theta$



Intuition Behind Cost Function


$$J(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \left(h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2$$



Intuition Behind Cost Function


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



Basic search Procedure

- Choose initial value for θ
- Until we reach a minimum:
 - Choose a new value for θ to reduce $J(\theta)$




Gradient Descent

- Initialize θ
- Repeat until convergence

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

simultaneous update
for $j = 0 \dots d$

learning rate (small)
e.g., $\alpha = 0.05$



- Initialize θ
- Repeat until convergence

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

simultaneous update
for $j = 0 \dots d$

For Linear Regression: $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)^2$

$$= \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right)^2$$

$$= \frac{1}{n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right) \times \frac{\partial}{\partial \theta_j} \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right)$$

$$= \frac{1}{n} \sum_{i=1}^n \left(\sum_{k=0}^d \theta_k x_k^{(i)} - y^{(i)} \right) x_j^{(i)}$$

- Initialize θ
- Repeat until convergence


$$\theta_j \leftarrow \theta_j - \alpha \frac{1}{n} \sum_{i=1}^n \left(h_{\theta} \left(\mathbf{x}^{(i)} \right) - y^{(i)} \right) x_j^{(i)} \quad \begin{array}{l} \text{simultaneous} \\ \text{update} \\ \text{for } j = 0 \dots d \end{array}$$

Assume convergence when $\|\boldsymbol{\theta}_{new} - \boldsymbol{\theta}_{old}\|_2 < \epsilon$

Gradient Descent


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Choosing α

α too small



slow convergence

α too large



Increasing value for $J(\theta)$

- May overshoot the minimum
- May fail to converge
- May even diverge

To see if gradient descent is working, print out $J(\theta)$ each iteration

- The value should decrease at each iteration
- If it doesn't, adjust α